Growing simplified vine copula trees: improving Dißmann's algorithm

Daniel Kraus[*] and Claudia Czado[†]

March 16, 2017

Abstract

Vine copulas are pair-copula constructions enabling multivariate dependence modeling in terms of bivariate building blocks. One of the main tasks of fitting a vine copula is the selection of a suitable tree structure. For this the prevalent method is a heuristic called Dißmann's algorithm. It sequentially constructs the vine's trees by maximizing dependence at each tree level, where dependence is measured in terms of absolute Kendall's τ. However, the algorithm disregards any implications of the tree structure on the simplifying assumption that is usually made for vine copulas to keep inference tractable. We develop two new algorithms that select tree structures focused on producing simplified vine copulas for which the simplifying assumption is violated as little as possible. For this we make use of a recently developed statistical test of the simplifying assumption. In a simulation study we show that our proposed methods outperform the benchmark given by Dißmann's algorithm by a great margin. Several real data applications emphasize their practical relevance.

Keywords: Vine copulas, structure selection, simplifying assumption, Dißmann's algorithm, test for constant conditional correlation.

1 Introduction

The list of references in which vine copulas are used to describe the dependence structure of multivariate random variables is long and does not cease to grow. After their inauguration in Aas et al. (2009) these pair-copula constructions have been applied to statistical problems in biology (Barthel et al., 2016; Schellhase and Spanhel, 2017), sociology (Cooke et al., 2015), hydrology (Erhardt and Czado, 2015; Hobæk Haff et al., 2015; Killiches and Czado, 2015; Pereira et al., 2016) and finance (Maya et al., 2015; Almeida et al., 2016; Kraus and Czado, 2017), to only name a few recent ones. An extensive review of vine copulas and their applications in finance is given in a survey by Aas (2016). In order to illustrate the idea of vine copulas we present an exemplary vine copula decomposition of a three-dimensional vine copula:

$$c(u_1, u_2, u_3) = c_{12}(u_1, u_2) \, c_{23}(u_2, u_3) c_{13;2} \left(C_{1|2}(u_1|u_2), C_{3|2}(u_3|u_2); u_2 \right). \qquad (1.1)$$

The multivariate copula density is expressed in terms of bivariate building blocks called pair-copulas. Two of them are unconditional (c_{12} and c_{23}) and one is conditioned ($c_{13;2}$

---

[*]Corresponding author, Zentrum Mathematik, Technische Universität München, Boltzmanstraße 3, 85748 Garching (email: daniel.kraus@tum.de)

[†]Zentrum Mathematik, Technische Universität München, Boltzmanstraße 3, 85748 Garching

on u_2). In general, there are various ways of specifying a vine copula decomposition. E.g. in the three-dimensional case, there are three possibilities depending on which variable is chosen to be the conditioning one. When we speak of the structure of the vine, we mean the specification which pairs of variables are modeled conditioned on which other variables. It is now a stylized fact that the flexibility of vines is able to overcome the problem of a limited model choice in the case when only elliptical or Archimedean copulas are used. Being able to assign a different bivariate copula to any pair of variables, the new issue arises of having too many possible models to choose from. Not only the selection of pair-copula families and corresponding parameters, but especially the vast number of possible model structures make the search of an "optimal" vine for a given data set a challenging task. Until now, almost unanimously, the method of choice for structure selection is the so-called Dißmann algorithm (Dißmann et al., 2013). This is also due to its prominent implementation as function `RVineStructureSelect` in the R-package `VineCopula` (Schepsmeier et al., 2016). Dißmann's algorithm is a heuristic that sequentially constructs the vine's structure, trying to capture most of the dependence in the lower trees.

When fitting a vine copula we usually make the so-called *simplifying assumption*, i.e. that the bivariate copulas of conditioned pairs of variables in the vine copula model do not depend on the value of the variable(s) they are conditioned on. In the three-dimensional example of Equation (1.1) this would imply that the conditional copula density $c_{13;2}$ would not depend on u_2. The implications and possible relaxation of this assumption have been discussed thoroughly in the recent literature (Killiches et al., 2017; Vatter and Nagler, 2016; Spanhel and Kurz, 2015; Acar et al., 2012, and many others). Even though the simplifying assumption is a model restriction, it is the only way to make inference for vine copulas tractable in higher dimensions.

The motivation of this paper is the observation that the model structure of a vine copula has a considerable influence on the validity of the simplifying assumption. To illustrate this, reconsider the three-dimensional vine copula model from Equation (1.1), now in its simplified form:

$$c(u_1, u_2, u_3) = c_{12}(u_1, u_2)\, c_{23}(u_2, u_3) c_{13;2}\left(C_{1|2}(u_1|u_2), C_{3|2}(u_3|u_2) \right).$$

This density can also be written as another vine copula using a different tree structure, e.g. the one containing the pairs $(1, 2)$ and $(1, 3)$ in the first and $(2, 3; 1)$ in the second tree:

$$c(u_1, u_2, u_3) = c_{12}(u_1, u_2)\, c_{13}(u_1, u_3) c_{23;1}\left(C_{2|1}(u_2|u_1), C_{3|1}(u_3|u_1); u_1 \right).$$

This decomposition does not need to be of simplified form anymore, i.e. the copula $c_{23;1}$ might now depend on the conditioning value u_1. Hence, given data from the simplified vine copula model we expect a fitted simplified vine copula with the original tree structure to have a significantly better model fit than one with a different tree structure. While Dißmann's structure selection algorithm only takes into account the strength of dependencies between pairs of variables, no consideration about the resulting conditional pair-copula is made. In this paper, we propose two new structure selection algorithms that try to amend exactly this flaw. By using a test for constant conditional correlations developed in Kurz and Spanhel (2017), we incorporate information about the violation of the simplifying assumption when constructing the vine's tree structure. The first algorithm selects the first tree similar to Dißmann's algorithm and uses the p-values of these tests as weights in the subsequent trees. The second algorithm we propose fits a C-vine where each root node is selected in a way that minimizes the occurrence of non-simplifiedness in the next tree. As we will see in a simulation study, this improves the model fit in terms of the Akaike information criterion (AIC) compared to Dißmann most of the time, especially when the

dimension is large. Finally, in revisiting several classic data sets that have already been studied in the vine copula context we demonstrate the practical relevance of our proposed methods.

The remainder of the paper is organized as follows. Section 2 provides the theoretical background needed to understand the concepts introduced in this paper. Two new algorithms for a vine's structure selection are motivated and proposed in Section 3. The performance of these algorithms is evaluated in a simulation study in Section 4. Real data applications are given in Section 5 and Section 6 concludes.

2 Theoretical Background

2.1 Vine copulas

We shortly recall the most important definitions needed for the construction of vine copulas. Details can be found in Bedford and Cooke (2002) or Aas et al. (2009). Moreover, we assume the concept of copulas to be known (refer to Nelsen (2007) and Joe (1997) for a thorough introduction to the topic) and restrict all our analyses to variables and data on the uniform copula scale $[0, 1]^d$.

A d-dimensional vine copula is a pair-copula construction consisting of $d(d-1)/2$ unconditional and conditional bivariate copulas, whose structure is organized by a set of linked trees $\mathcal{V} = (T_1, \ldots, T_{d-1})$ satisfying

(i) $T_1 = (V_1, E_1)$ is a tree with nodes $V_1 = \{1, \ldots, d\}$ and edges E_1. A tree is understood as a graph where any two nodes are connected by a unique path (refer to Diestel, 2005, for an introduction to graph theory).

(ii) For $m = 2, \ldots, d-1$, the tree T_m consists of nodes $V_m = E_{m-1}$ and edges E_m.

(iii) For $m = 2, \ldots, d-1$, two nodes of T_m can only be connected by an edge if the corresponding edges of T_{m-1} have a common node.

Each edge e of the vine copula model's $d-1$ trees is associated with a bivariate pair-copula $c_{j_e,k_e;D_e}$, where – following the notation of Czado (2010) – j_e and k_e denote the indices of the conditioned variables U_{j_e} and U_{k_e}, and D_e represents the conditioning set corresponding to edge e. Thus, $c_{j_e,k_e;D_e}$ is the density of the copula between random variables $U_{j_e|D_e}$ and $U_{k_e|D_e}$, where $U_{i|D} := C_{i|D}(U_i|U_D)$. The vine density can then be written as

$$c(u_1, \ldots, u_d) = \prod_{m=1}^{d-1} \prod_{e \in E_m} c_{j_e k_e;D_e} \left(C_{j_e|D_e}(u_{j_e}|\mathbf{u}_{D_e}), C_{k_e|D_e}(u_{k_e}|\mathbf{u}_{D_e}); \mathbf{u}_{D_e} \right). \qquad (2.2)$$

Condition (iii) for the set of trees $\mathcal{V}$ is also known as the *proximity condition*. When one constructs the trees sequentially, starting with T_1, then it induces constraints on the nodes in the next tree which can be connected by an edge. Two types of trees play a special role regarding the proximity condition. For $m = 1, \ldots, d-2$, if tree T_m has a star like structure, i.e. there is one node j connected to all the other nodes, then the proximity condition imposes no restriction for the construction of tree T_{m+1}, since all nodes share the common node j from T_m. Thus, all nodes of tree T_{m+1} are allowed to be connected. On the contrary, if tree T_m has a path like structure, i.e. each node has at most two neighbors, then all higher trees $T_{m+1}, \ldots, T_{d-1}$ are already determined by the proximity condition, since for every node in T_{m+1} there exist at most two nodes that can be connected to it, resulting in a path like structure for tree T_{m+1} as well. A vine structure where all trees consist of paths is called a *D-vine*, while a vine with only star like trees is known as a *C-vine*.

2.2 Dißmann's algorithm for selecting simplified vine copulas

Since the sequential structure selection algorithm proposed in Dißmann et al. (2013) will serve as a benchmark vine model selection method, we shortly outline how it works. Assume we are given an i.i.d. sample of size n of a d-dimensional copula, denoted by $(u_j^{(i)})_{j=1,\ldots,d}^{i=1,\ldots,n}$. Since a vine copula model consists of three parts (the vine structure, the parametric families of each pair-copula and their corresponding copula parameters), each of these components have to be estimated.

Assume at first that the tree structure is given. Then for each of the $d(d-1)/2$ pair-copulas a family and copula parameter have to be selected. Starting with the first tree, for each unconditional pair the maximum likelihood estimate of the copula parameter is determined for each family studied (see Appendix A for a list of families currently implemented in the `VineCopula` package). Then the family with the lowest Akaike information criterion (AIC) is chosen (Akaike, 1998). In the higher trees, for each edge corresponding to a conditional pair-copula $c_{j_e,k_e;D_e}$, the estimated pseudo observations $u_{j_e|D_e}^{(i)} := \hat{C}_{j_e|D_e}(u_{j_e}^{(i)}|\mathbf{u}_{D_e}^{(i)})$ and $u_{k_e|D_e}^{(i)} := \hat{C}_{k_e|D_e}(u_{k_e}^{(i)}|\mathbf{u}_{D_e}^{(i)})$, $i = 1, \ldots, n$, are calculated and used to find the optimal family and parameter for pair-copula $c_{j_e,k_e;D_e}$. This is repeated until all pair-copulas are fitted. Note that it is possible to incorporate a test for independence in the pair-copula selection process (see Genest and Favre, 2007). Based on the estimated Kendall's τ the null hypothesis that the (pseudo) observations come from the independence copula is tested. Then, if the null hypothesis is not rejected for some chosen level β, the pair-copula to be estimated is chosen to be the independence copula.

Regarding the selection of the tree structure, Dißmann's algorithm uses a heuristic that models the strongest pairwise dependencies (measured in terms of the absolute empirical Kendall's τ value) in the lower trees. To be precise, the algorithm starts with the first tree, where all pair-wise empirical Kendall's τ values are determined and then a maximum spanning tree using the absolute τ values as weights is selected (e.g. by the Algorithm of Prim (Cormen et al., 2001, Section 23.2)). For the next tree, all required pseudo observations are determined as above and for all edges allowed by the proximity condition the Kendall's τ values are estimated. Again, a maximum spanning tree is selected using these empirical Kendall's τ values as edge weights, now on the graph constrained by the proximity condition. This procedure is iterated until all $d-1$ trees are selected.

Dißmann et al. (2013) justify this heuristic by noting that the lowest trees have the greatest influence on the overall fit and thus it is important to model most of the dependence early. Further, they argue that this procedure minimizes estimation errors in higher trees, because typically the dependence decreases when using the algorithm, making rounding errors less severe. As a measure of the strength of dependence Kendall's τ is used, since it is a rank-based dependence measure facilitating the comparison of different copula families (see e.g. Nelsen, 2007, Ch. 5.1.1).

Alternative methods using other edge weights for the determination of the maximum spanning tree have been proposed. For example, the fitted pair-copula's AIC or the p-value of a goodness-of-fit test have been used in Czado et al. (2013), but they are not commonly used, since all possible pair-copulas have to be estimated first to obtain all the edge weights. This becomes time consuming especially in higher dimensions. Another approach for the tree selection of vines is given by Kurowicka (2011), who also uses the heuristic of modeling most of the dependence in the lowest trees by starting with the last tree, where the edge with the lowest partial correlation is chosen. This is repeated, going backwards, until all trees are specified.

So we see that there is already quite a range of structure selection methods. However,

all these approaches do not take into account the selected structure's implications on the validity of the simplifying assumption and therefore on the overall model fit. The methods we will propose in Section 3 remedy this issue. They are based on a test whether the simplifying assumption is fulfilled for a considered pair-copula term as described in the following section.

2.3 Test for constant conditional correlations (CCC test)

Formally, the simplifying assumption requires that all the pair-copulas of the vine appearing in Equation (2.2) can be written as follows:

$$c_{j_e k_e; D_e}\left(\cdot, \cdot; \mathbf{u}_{D_e}\right) \equiv c_{j_e k_e; D_e}\left(\cdot, \cdot\right).$$

This means that the pair-copula is independent of the conditioning variables, implying that the dependence associated with the copula is constant with respect to $\mathbf{u}_{D_e}$. A stochastic representation of the simplifying assumption is given by $(U_{j_e|D_e}, U_{k_e|D_e}) \perp \mathbf{U}_{D_e}$, denoting that the random variables $U_{j_e|D_e}$ and $U_{k_e|D_e}$ are jointly independent of $\mathbf{U}_{D_e}$ (Kurz and Spanhel, 2017). Without the simplifying assumption, one would have to allow $c_{j_e k_e; D_e}$ to depend on $\mathbf{u}_{D_e}$, making it a $(|D_e| + 2)$-dimensional function to be estimated (which is d-dimensional for the pair-copula in the last tree). This would defeat the whole purpose and convenience of vine models, whose idea it is to express a d-dimensional copula just in terms of *bivariate* building blocks. While for low dimensions of the conditioning vector researchers have found ways to relax the simplifying assumption (Acar et al., 2012; Vatter and Nagler, 2016; Schellhase and Spanhel, 2017), in higher dimensions this does not seem feasible. Further discussion and implications of the simplifying assumption can be found in (Haff et al., 2010; Stöber et al., 2013; Spanhel and Kurz, 2015; Killiches et al., 2016). Recently, Kurz and Spanhel (2017) developed a statistical testing method assessing the severity of the violation of the simplifying assumption for each conditional pair-copula in a vine copula. It tests the null hypothesis that the conditional correlation $\rho_{j_e k_e|D_e}$ associated with pair-copula $C_{j_e k_e; D_e}$ is constant with respect to the conditioning variables $\mathbf{U}_{D_e}$. For this purpose, the support Ω_0 of $\mathbf{U}_{D_e}$ is divided by a partition $\Gamma := \{\Omega_1, \dots, \Omega_L\}$, with $L \in \mathbb{N}$, $\Omega_i \cap \Omega_j = \emptyset$, for $i \neq j$, and $P(\mathbf{U}_{D_e} \in \Omega_i) > 0$, for $i = 1, \dots, L$. The idea of the test is that under the null hypothesis the conditional correlation between $U_{j_e|D_e}$ and $U_{k_e|D_e}$ given $\mathbf{U}_{D_e} \in \Omega_i$ does not depend on i. Hence, denoting $\rho_{\Omega_i} := Corr(U_{j_e|D_e}, U_{k_e|D_e}|\mathbf{U}_{D_e} \in \Omega_i)$, the null hypothesis is equivalent to testing $\rho_{\Omega_1} = \dots = \rho_{\Omega_L}$, which we will call the constant conditional correlation (CCC) assumption. To derive a test statistic we observe that for the vector of estimated conditional correlations $\hat{R}^{(n)}(\Gamma) := (\hat{\rho}_{\Omega_1}^{(n)}, \dots, \hat{\rho}_{\Omega_L}^{(n)})'$, estimated by the sample Pearson's correlation coefficient based on an i.i.d. sample of size n, it holds

$$\sqrt{n}(\hat{R}^{(n)}(\Gamma) - R(\Gamma)) \xrightarrow[n \to \infty]{d} N(\mathbf{0}, \Sigma(\Gamma)),$$

where $R(\Gamma)$ is the vector of true conditional correlations and $\Sigma(\Gamma)$ is the asymptotic variance-covariance matrix (see Kurz and Spanhel, 2017). The asymptotic normality of $\hat{R}^{(n)}(\Gamma)$ is used to derive an asymptotically χ^2 distributed test statistic by taking differences and normalizing. Defining the matrix $A \in \{-1, 0, 1\}^{L-1 \times L}$ with $A_{ij} = \mathbb{1}\{i = j\} - \mathbb{1}\{i = j - 1\}$, $i = 1, \dots L - 1$, $j = 1, \dots L$, i.e.,

$$A = \begin{pmatrix} 1 & -1 & & & \\ & 1 & -1 & & \\ & & \ddots & \ddots & \\ & & & 1 & -1 \end{pmatrix},$$

where the omitted entries equal zero, we have that

$$A\hat{R}^{(n)}(\Gamma) = \begin{pmatrix} \hat{\rho}_{\Omega_1}^{(n)} - \hat{\rho}_{\Omega_2}^{(n)} \\ \hat{\rho}_{\Omega_2}^{(n)} - \hat{\rho}_{\Omega_3}^{(n)} \\ \vdots \\ \hat{\rho}_{\Omega_{L-1}}^{(n)} - \hat{\rho}_{\Omega_L}^{(n)} \end{pmatrix}.$$

The real-valued test statistic is defined by

$$T_n(\Gamma) = (A\hat{R}^{(n)}(\Gamma))'(A\hat{\Sigma}^{(n)}(\Gamma)A')^{-1}A\hat{R}^{(n)}(\Gamma),$$

where $\hat{\Sigma}^{(n)}(\Gamma)$ is a consistent estimator for $\Sigma(\Gamma)$. Since Kurz and Spanhel (2017) show that under regularity conditions and H_0 it holds that

$$nT_n(\Gamma) \xrightarrow[n\to\infty]{d} \chi^2_{L-1},$$

an asymptotic β-level test is given by

$$\mathbb{1}\{nT_n(\Gamma) \geq F^{-1}_{\chi^2_{L-1}}(1-\beta)\},$$

where $F^{-1}_{\chi^2_{L-1}}(1-\beta)$ denotes the $1-\beta$ quantile of the χ^2 distribution with $L-1$ degrees of freedom.

Of course, the power of this test highly depends on the chosen partition Γ. Therefore, Kurz and Spanhel (2017) adjust the test statistic to accommodate for the combination of different partitions and further use decision trees to find partitions where a possible violation of the simplifying assumption is most pronounced. The partition size L is determined by the algorithm and grows with $|D_e|$ (see Kurz and Spanhel, 2017, for details). Finally, they show that the test has a very high observed power compared to benchmark methods. The test is implemented as function `pacotest` in the R package `pacotest` (Kurz, 2017).

3   Two new tree selection algorithms

3.1   Motivation

As a motivation we first consider three-dimensional data sets. Here we know that there are three different possible vine tree structures and that the specification of the first tree already fully describes the whole tree structure. In Section 2.2 we have seen that Dißmann's algorithm chooses the tree structure that maximizes the dependence between the variables in the first tree. However, the resulting tree structure might yield a non-simplified vine copula even though the true model is a simplified vine copula. Hence, the model fit of Dißmann in terms of log-likelihood might not be optimal.

This motivates us to select tree structure by testing which combination of variables is "most simplified". In detail, we use the R-function `pacotest` to test for the hypothesis that the CCC assumption holds for $U_i, U_j|U_k$ with $k \in \{1, 2, 3\}$ and $\{i, j\} = \{1, 2, 3\}\backslash k$. Then, we choose the tree structure which has the highest p-value in the second tree.

Example 1   We show how this works for a real life data set, namely a subset of the well-known seven-dimensional hydro-geochemical data set first investigated by Cook and Johnson (1986), consisting of $N = 655$ observations of log-concentrations of the three

chemicals cobalt (U_1), titanium (U_2) and scandium (U_3) in water samples taken from a river near Grand Junction, Colorado. This data set has been examined with regard to the simplifying assumption by many researchers, e.g. Acar et al. (2012), Killiches et al. (2017) and Killiches et al. (2016). After transforming the data to the copula scale by applying the empirical probability integral transform, we estimate the Kendall's τ values for each of the three pairs. They are $\hat{\tau}_{12} = 0.54$, $\hat{\tau}_{13} = 0.36$ and $\hat{\tau}_{23} = 0.44$. Consequently, Dißmann's algorithm chooses edges $(1, 2)$ and $(2, 3)$ corresponding to pair-copulas c_{12} and c_{23} for the first tree resulting in the conditional copula $c_{13;2}$ to be modeled in the second tree. The Dißmann algorithm fits all pair-copulas as t-copulas. The log-likelihood and AIC-values of the vine copula fitted by Dißmann's algorithm are 428.8 and -845.6, respectively (see also Structure 2 in Table 1).

Structure	Tree 1	conditional copula	p-value	sum of τ	log-lik	AIC
1	2–1–3	$c_{23;1}$	0.29	$0.54 + 0.36 = 0.90$	434.9	-857.8
2	1–2–3	$c_{13;2}$	0.01	$0.54 + 0.44 = 0.98$	428.8	-845.6
3	1–3–2	$c_{12;3}$	0.19	$0.36 + 0.44 = 0.80$	418.5	-825.1

Table 1: The p-values, sums of Kendall's τ values, log-likelihoods and AIC-values of the three possible tree structures for the three-dimensional uranium dataset.

Let us now examine the p-values of the test for constant conditional correlations for the three possible tree structures. They are 0.29, 0.01, 0.19 for the conditional copulas $c_{23;1}$, $c_{13;2}$ and $c_{12;3}$, respectively. Hence, we see that the structure chosen by Dißmann's algorithm has a strong indication for non-simplifiedness, such that we would rather choose one of the other two structures. The structure corresponding to the highest p-value (Structure 1 in Table 1) also has a larger sum of Kendall's τ values in the first tree compared to the other structure with a large p-value ($0.54 + 0.36 = 0.90$ vs. $0.36 + 0.44 = 0.80$), such that we would choose the structure with edges $(1, 2)$ and $(1, 3)$ in the first and $(2, 3; 1)$ in the second tree. The fitted vine copula of Structure 1 has a Tawn and a t-copula in the first tree and a BB7 copula in the second tree. With a log-likelihood of 434.9 and an AIC of -857.8 it provides a better model fit than the one selected by Dißmann's algorithm. For completeness we note that Structure 3 in Table 1 has the worst fit, with the lowest sum of τ in the first tree and rather low p-value for the conditional copula in the second tree. All in all, with Structure 1 we found a way of describing this data set using a simplified vine copula, whereas in the recent literature it seemed necessary to use a non-simplified vine copula model when Structure 2 was used.

Extending this idea to higher dimensions, things get more complicated because there are superexponentially more tree structures to choose from. While e.g. in five dimensions there are still "only" 480 possible tree structures, in ten dimensions they already explode to $4.8705 \cdot 10^{14}$. See Morales-Nápoles (2011) for a general formula for counting the number of different vine structures. The question is now how to incorporate the information about the tests concerning the simplifiedness of the conditional copulas for the determination of the best vine tree.

3.2 Algorithm 1: Regular vine structure selection using CCC test based weights

The first approach that we propose is to choose the first tree (where there are no explicit considerations about simplifiedness yet) according to Dißmann's approach and then to

incorporate the CCC tests in the higher level trees. In each tree, the test's p-values of all edges allowed by the proximity condition are calculated and used to find the maximum spanning tree. This would constitute a compromise between both approaches reflecting as much of the dependence as possible in the first tree and accounting for non-simplifiedness in the trees where conditional copulas have to be fitted. Further, one can also assign each edge a score combining both the edge's p-value and estimated absolute Kendall's τ. Hence, the score of an edge e would be defined as

$$s_\alpha(e) := \alpha \cdot f_p(e) + (1 - \alpha) \cdot f_\tau(e),$$

where α is a weighting factor and the functions f_p and f_τ map each edge e allowed by the proximity condition to a score regarding its degree of simplifiedness and absolute Kendall's τ values, respectively. For example, f_p could be the function mapping each edge to its associated rank of the p-value arising from the CCC test, i.e. $f_p(e_i) = i$, if the edges e_i are ordered such that p-value$(e_1) <$ p-value$(e_2) < \ldots <$ p-value(e_N). Here, N is the number of edges allowed by the proximity condition. Similarly, f_τ could rank the edges by their absolute empirical Kendall's τ values, such that the scales of f_p and f_τ would coincide. The weighting factor α can be chosen by the user. We will see in Section 4.2 that values around 0.6 yield good fits in terms of the AIC. After having determined the score of each edge, a maximum spanning tree on the edges using the scores $s_\alpha(e)$ is constructed. Then, similar to Dißmann's algorithm the pair-copulas corresponding to the chosen tree are fitted (with or without an independence test). Algorithm 1 implements this procedure.

Algorithm 1 Regular vine structure selection using CCC test based weights

---

Input: d-dimensional copula data, weighting factor α, edge score functions f_p and f_τ.

1: Tree 1: estimate absolute empirical Kendall's τ values between all variables and find the corresponding maximum spanning tree. For each selected edge of the tree, fit a pair-copula based on the copula data and determine the corresponding pseudo observations for the next tree.

2: **for** $m = 2, \ldots, d - 1$ **do**

3:     Tree m: for all edges e allowed by the proximity condition, calculate the score $s_\alpha(e) = \alpha \cdot f_p(e) + (1-\alpha) \cdot f_\tau(e)$ using the pseudo observations and find the corresponding maximum spanning tree. For each selected edge of the tree, fit a pair-copula based on the pseudo observations and determine the corresponding pseudo observations for the next tree.

Output: R-vine with tree structure focused on constant conditional dependencies as well as large dependencies in lower trees.

Example 2 We show how the algorithm works in detail for a 5-dimensional example. At first we generate a random R-vine object with random tree structure, families and parameters and simulate a sample of size 1000 from it (exact details on how the random vine is generated can be found in Section 4). The resulting true vine specification is a D-vine with pair-copulas given in Table 2.

The chosen vine's first tree connects the edges $(2, 3)$, $(2, 4)$, $(1, 5)$ and $(4, 5)$. In a first step we fit an R-vine to the simulated data using the function `RVineStructureSelect` in which Dißmann's algorithm is implemented. The resulting first two trees are displayed in Figure 1 as the gray graphs.

Tree 1	Tree 2
C_{23}: Clayton (0.44) ($\tau_{23} = 0.18$)	$C_{34;2}$: survival BB6 (1.07,2.31) ($\tau_{34;2} = 0.58$)
C_{24}: type 1 Tawn (2.43, 0.81) ($\tau_{24} = 0.51$)	$C_{25;4}$: BB8 (1.62, 0.7) ($\tau_{25;4} = 0.1$)
C_{15}: type 2 Tawn (6, 0.43) ($\tau_{15} = 0.39$)	$C_{14;5}$: 270 degree rotated type 1 Tawn (-3.89,
C_{45}: survival Gumbel (2.69) ($\tau_{45} = 0.63$)	0.54) ($\tau_{14;5} = -0.45$)
Tree 3	Tree 4
$C_{35;24}$: survival Joe (2.31) ($\tau_{35;24} = 0.42$)	$C_{13;245}$: survival BB1 (0, 2.39) ($\tau_{13;245} = 0.58$)
$C_{12;45}$: Frank (18.67) ($\tau_{12;45} = 0.8$)	

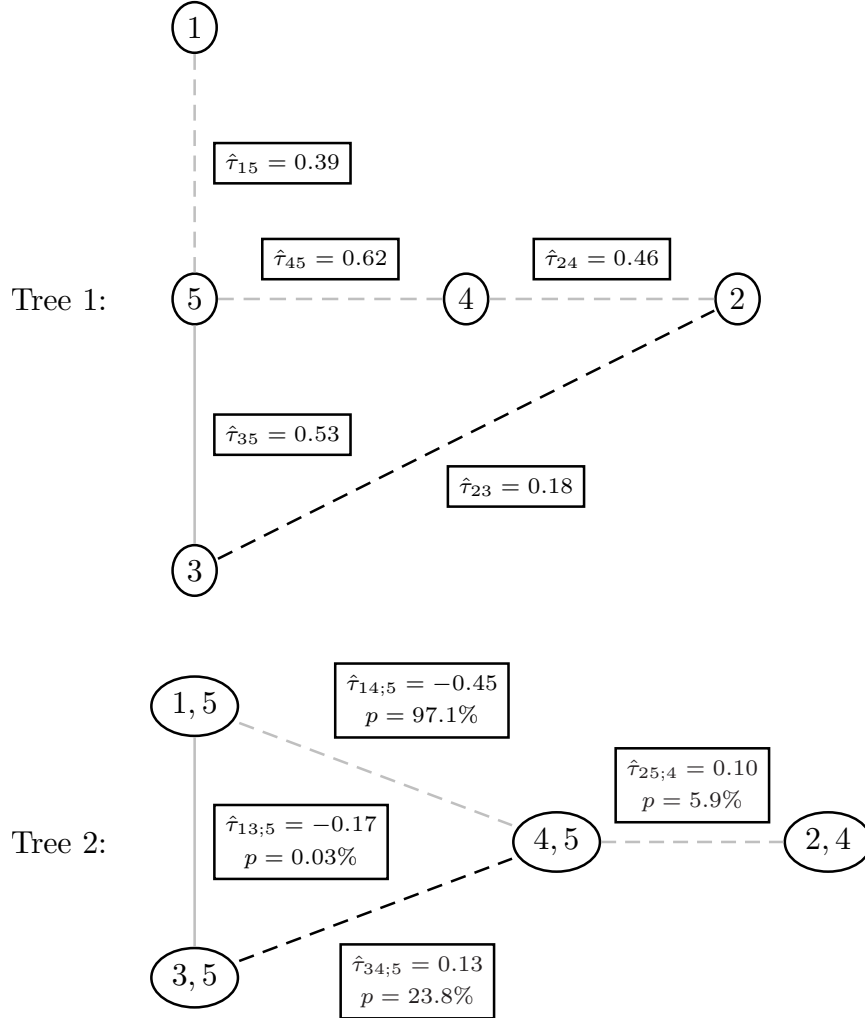Table 2: D-vine copula specification of Example 2.

Figure 1: First two trees selected by Dißmann's algorithm (gray). In Tree 1, the empirical Kendall's τ values of each edge are given in the boxes and the dashed lines correspond to the true tree structure. In Tree 2, additionally to Kendall's τ, the p-values of the conditional correlation test are given and the dashed lines mark the tree selected by Algorithm 1.

The first tree of the fitted R-vine is quite close to the true one (highlighted by a dashed line). Edges $(2, 4)$, $(1, 5)$ and $(4, 5)$ are selected and the true copula families are chosen with parameters very close to the true ones ($\hat{\tau}_{24} = 0.46$, $\hat{\tau}_{15} = 0.39$ and $\hat{\tau}_{45} = 0.62$). Only edge $(3, 5)$ is selected instead of $(2, 3)$ since it has a higher empirical Kendall's τ (0.53 compared to 0.18).

Due to the proximity condition in the second tree, node $(2, 4)$ can only be connected to

$(4, 5)$ (since it is the only other node containing a 4 or a 2), resulting in the edge $(2, 5; 4)$. Regarding nodes $(1, 5)$, $(3, 5)$ and $(4, 5)$ the proximity condition allows all of these nodes to be connected. Only focusing on the absolute empirical Kendall's τ value, Dißmann's algorithm chooses edges $(1, 3; 5)$ and $(1, 4; 5)$, since $|\hat{\tau}_{14;5}| > |\hat{\tau}_{13;5}| > |\hat{\tau}_{34;5}|$. Altogether, the selected second tree constitutes a path and therefore the higher level trees are already determined by the proximity condition. The AIC of the R-vine fitted by Dißmann's algorithm is -7775.5.

Next, we consider the results of our proposed algorithm on this data set. Of course, the selected first tree coincides with the one from Dißmann's algorithm. However, in the second tree, when choosing which of the nodes $(1, 5)$, $(3, 5)$ and $(4, 5)$ to connect, we see that even though edge $(1, 3; 5)$ has a slightly larger absolute Kendall's $\hat{\tau}$ than edge $(3, 4; 5)$, its p-value (0.03%) is much smaller than the 23.8% of edge $(3, 4; 5)$. Thus, for $\alpha \geq 0.5$ our proposed algorithm would choose edges $(3, 4; 5)$ and $(1, 4; 5)$ for the second tree (cf. the dashed graph in the lower panel of Figure 1). After having fit the third and fourth trees accordingly, the overall resulting R-vine proves to yield a much better fit than Dißmann's algorithm with an AIC of -8272.9.

In the simulation study presented in Section 4 we repeat this procedure 1000 times and find that Algorithm 1 achieves a better or equal AIC-value in 80.5 percent of the simulations.

3.3 Algorithm 2: C-vine structure selection using CCC test based weights

One major disadvantage of Algorithm 1 is that the overall tree structure is already strongly affected by the first tree due to the proximity condition. For example, if the first tree is fitted with a path-like structure, then the whole vine is automatically a D-vine since the proximity condition only allows for one possible (also path-like) structure in the higher trees. In this case we cannot account for a possible non-simplifiedness of the model at all. For example, in three dimensions Algorithm 1 would always yield the same structure as the Dißmann algorithm since for three nodes every tree constitutes a path. So, e.g. for the three-dimensional uranium example described in Section 3.1 where the consideration of the test for the simplifying assumption improved the structure fit, the proposed algorithm would be futile.

Thus we have to find a way how to sequentially select a vine structure incorporating in the current tree the information about which conditional correlations are not constant in the subsequent tree while still keeping the subsequent graph as flexible as possible. We have seen that a path-like structure in the first tree already determines the structure of the remaining trees. Conversely, we know that we retain the highest flexibility by fitting a star-like structure in the first tree since then the proximity condition imposes no restriction (such that all nodes in the subsequent tree are allowed to be linked by an edge). Following this logic for all trees, a star-like structure is reasonable for every tree level. The root node v_m of each star should be chosen that on the one hand the sum of absolute empirical Kendall's τ values between this root node and all the other variables (conditioned on the earlier root nodes $v_1, \ldots, v_{m-1}$) is large. On the other hand, conditioned on the root node v_m (and all the earlier root nodes) pairs of the remaining variables should have rather constant conditional correlations, i.e. the p-values of the CCC test conditioning on the root nodes should be large.

Formally, at tree level m we assign each remaining node $v \in N_m := \{1, \ldots, d\} \backslash \{v_1, \ldots, v_{m-1}\}$ the score $s_\alpha(v)$ (operating on nodes v instead of edges e as in Algorithm 1) defined as

$$s_\alpha(v) = \alpha \cdot g_p(v) + (1 - \alpha) \cdot g_\tau(v).$$

10

We then choose the one with the highest score as a root node for tree m. Again, α is a weighting factor and g_p and g_τ are functions mapping a node to ranked scores regarding the CCC tests and the Kendall's τ values, respectively. In particular, g_p calculates for all $v \in N_m$ the p-value score

$$p(v) := \sum_{i,j \in N_m \setminus v, \, i \neq j} r(p_{ij;v_1,\ldots,v_{m-1},v}),$$

and then maps each v to its rank among all p-value scores $\{p(v)\}_{v \in N_m}$ (e.g. the root node with the smallest p-value score would be assigned rank 1). Here, $p_{ij;v_1,\ldots,v_{m-1},v}$ denotes the p-value of the test for constant conditional correlation of $C_{ij;v_1,\ldots,v_{m-1},v}$ and the function r maps a p-value $p_{i_0j_0;v_1,\ldots,v_{m-1},v_0}$ to its rank among all possible p-values $p_{ij;v_1,\ldots,v_{m-1},v}$ with $i, j, v \in N_m$ pairwise distinct. Instead of choosing the function r to be the rank transformation, other transformations such as the logarithm are possible. However, the results of the simulation study show that the rank transformation yields the best results (see Appendix F). Regarding Kendall's τ, g_τ similarly calculates the τ scores

$$t(v) := \sum_{i \in N_m \setminus v} |\hat{\tau}_{iv;v_1,\ldots,v_{m-1}}|, \; v \in N_m$$

and then assigns each v to its corresponding rank among these τ scores. The Kendall's τ associated to $c_{iv;v_1,\ldots,v_{m-1}}$ is estimated as the empirical Kendall's τ between the pseudo observations $\hat{\mathbf{u}}_{i|v_1,\ldots,v_{m-1}}$ and $\hat{\mathbf{u}}_{v|v_1,\ldots,v_{m-1}}$ and is denoted by $\hat{\tau}_{iv;v_1,\ldots,v_{m-1}}$.

After the optimal root node is found, similar to the other algorithms all pair-copulas implied by the resulting tree have to be fitted (with or without independence test) for the calculation of the pseudo observations of the next tree level (see Algorithm 2).

Algorithm 2 C-vine structure selection using CCC test based weights
---
Input: d-dimensional copula data, weighting factor α, node score functions g_p and g_τ.
 1: **for** $m = 1, \ldots, d-1$ **do**
 2:    Tree m: Determine the optimal root node by evaluating score $s_\alpha(v) = \alpha \cdot g_p(v) + (1 - \alpha) \cdot g_\tau(v)$ for all nodes $v \in N_m$ and choose the maximal one. Fit pair-copulas between this root node and the remaining nodes based on the pseudo observations. Calculate the corresponding pseudo observations needed for the tree selection in the next step.

Output: C-vine with tree structure focused on constant conditional dependencies as well as large dependencies in lower trees.

Using this algorithm with any $\alpha > 0.5$, we retrieve the AIC-optimal tree structure for the uranium example discussed in Section 3.1.

Example 2 (continued) We revisit the 5-dimensional example from the previous section and want to apply Algorithm 2 to the simulated data (with $\alpha = 0.6$ and rank-transformed p-values). Table 3 displays the information needed to find the optimal root node v_1 in the first tree of the C-vine.

For example, the p-value score $p(1) = 91$ is the sum of the ranks of the 6 p-values $p_{ij;1}$, $i, j \in \{2, 3, 4, 5\}$, $i \neq j$, among the 30 possible p-values. Since 91 is the third largest p-value score, its rank among the p-value scores is 3. Similarly, since $\sum_{i=2}^{5} |\hat{\tau}_{i1}| = 1.04$ is the smallest of the τ scores node 1 gets assigned rank 1. Hence, the overall score is calculated as $s_\alpha(v) = 0.6 \cdot 3 + 0.4 \cdot 1 = 2.2$. Doing this for the other possible root nodes we see that node 5 has the largest p-value and τ scores and therefore with an overall score of 5.0 is selected as

node v	$p(v)$	rank($p(v)$)	$t(v)$	rank($t(v)$)	$s_\alpha(v)$
1	91	3	1.04	1	2.2
2	89	2	1.47	3	2.4
3	105	4	1.31	2	3.2
4	69	1	1.72	4	2.2
5	111	5	1.98	5	5.0

Table 3: p-value scores $p(v)$ and τ scores $t(v)$ with associated ranks, and the overall score $s_\alpha(v)$ with $\alpha = 0.6$ for the determination of the root node v_1 in the first tree of the C-vine fitted by Algorithm 2.

the first root node. This procedure is repeated for all trees yielding a C-vine with root node ordering 5–4–2–1–3 and an AIC of -7446.6. While we see that this model has a higher AIC than the ones selected by Algorithm 1 and by Dißmann's algorithm (which may be due to the fact that the true model is a D-vine and Algorithm 2 fits a C-vine), it still finds the AIC-optimal out of the 60 possible C-Vines structures. Further, we will see in Section 4, where this procedure is repeated 1000 times, that in 73.6% of the time Algorithm 2 finds a C-vine with a better AIC than the one of Dißmann's chosen R-vine.

4 Simulation study

4.1 Setup

We perform an extensive simulation study, evaluating the performance of the two proposed algorithms in many different scenarios. For each scenario, we repeat $R = 1000$ simulations of sample size n of randomly generated d-dimensional R-vine copulas. For these, we sample uniformly one out of the $2^{(d-2)(d-3)/2}$ different structure matrices with natural ordering as described in Joe et al. (2011) using the R function `RVineMatrixSample` from `VineCopula`; for each pair-copula, a random copula family out of the 12 families currently implemented in the package `VineCopula` is selected (see Appendix A for a list of the implemented families). Regarding the copulas' parameters, $Beta(2,2)$ distributed Kendall's τ values are generated and multiplied by -1 with a probability of 50% (note that the densities of some families have to be rotated by 90 degrees to accommodate negative dependence). The $Beta(2,2)$ distribution is symmetric around its mode 0.5 with a variance of 0.05 and has a 95% confidence interval given by $[0.09, 0.91]$. Next, for the two-parametric copula families the second parameter is randomly generated from suitable distributions (see Appendix B for details). Finally, the first parameter of all pair-copulas is derived from the sampled Kendall's τ values and the second parameters (where required) using R function `BiCopTau2Par` (see Joe, 1997, for details). An exemplary random 5-dimensional R-vine is given in the Table 2 of Section 3.2.

In the simulation study, for each of the samples of size n from these random R-vines we will calculate the AIC-values of the vines fitted using our proposed algorithms (here without the mentioned independence tests). The AIC-values will be compared to those of the fitted Dißmann vines. Before presenting the results of the simulation study, we shortly investigate the influence of the weighting factor α on the AIC-values of the models fitted by our algorithms.

4.2 Choice of weighting factor α

In the simulation study we will consistently use the weighting factor $\alpha = 0.6$. This choice is justified because we heuristically observe that the AIC of our fitted values on average is lowest for medium sized weighting factors with the best performances for $\alpha = 0.6$. We generate 1000 random d-dimensional vine copulas (as described in Section 4.1), simulate a sample of size n from each copula and apply our algorithms to the sample using weighting factors $\alpha \in \{0, 0.1, 0.2, \ldots, 0.9, 1\}$. For each α we compute the AIC of the fitted models averaged over the 1000 repetitions. Overall, we find that in most scenarios $\alpha = 0.6$ yields the lowest average AIC or close to it. Figure 2 displays exemplary results for Algorithm 2 in the setting $n = 1000$ and $d = 5, 10, 30$.

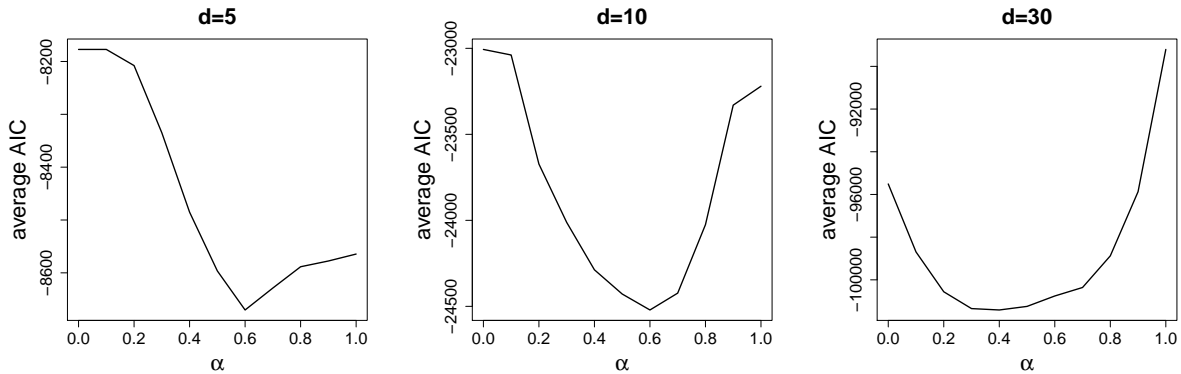
Figure 2: AIC-values of vine copulas fitted by Algorithm 2 depending on the weighting factor α in the setting $n = 1000$ and $d = 5, 10, 30$, averaged over 1000 repetitions.

We observe a U-shape of the plots indicating that choosing the weighting factor too small or too large does not yield optimal fits. So, focusing exclusively on the Kendall's τ values or the p-values of the CCC test apparently is not sufficient. For $d = 5$ and $d = 10$ the minimum of the average AIC-values is attained for $\alpha = 0.6$ and in dimension thirty for $\alpha = 0.4$ with the average AIC of $\alpha = 0.6$ being still quite close to the minimum. The results for sample sizes $n = 400$ and $n = 3000$ given in Appendix C are very similar and lead the same conclusion of $\alpha = 0.6$ being optimal. Finally, the same study applied for Algorithm 1 also yields an optimum of $\alpha = 0.6$.

4.3 Results

In our simulation study we let the dimension d of the randomly generated vine copulas vary between 3 and 50 and the sizes of the simulated samples n between 400 and 3000. The percentage of times where our algorithms performs better or equal than Dißmann is given in Table 4 (the percentages of equal performance are given in brackets and are omitted if they are 0). The results of the table are also visualized in Figure 6 of Appendix G.
We see that in general our two proposed algorithms perform very well compared to Dißmann's algorithm. Of course, in low dimensions Algorithm 1 performs very similar to Dißmann's algorithm since there the first tree greatly determines the model fit. So as we already noted, in three dimensions Algorithm 1 and Dißmann's algorithm always find the same vine structure and even in four and five dimensions they coincide in almost 90% and 70% of the times, respectively. In higher dimensions, when the selected tree structures differ more often, Algorithm 1 manages to find a better structure than Dißmann's algorithm in more than one third of the simulations for dimensions 10-15 and around one fourth of the

		d								
Algorithm	n	3	4	5	7	10	15	20	30	50
	400	100 (100)	92.2 (89.0)	80.8 (69.3)	54.9 (30.8)	41.9 (5.9)	34.8 (0.2)	30.0	26.5	22.1
Algorithm 1	1000	100 (100)	93.1 (89.0)	82.3 (69.7)	59.9 (32.3)	41.4 (4.8)	35.9 (0.1)	30.3	28.9	27.3
	3000	100 (100)	92.6 (88.6)	79.9 (68.0)	58.3 (31.3)	41.6 (4.2)	36.2 (0.2)	35.3	21.8	18.2
	400	84.4 (49.6)	68.8 (3.1)	67.7	75.9	82.0	93.2	96.7	99.2	99.8
Algorithm 2	1000	91.1 (51.5)	75.5 (3.8)	73.6	80.2	88.7	95.8	99.0	99.6	100
	3000	93.4 (47.6)	77.1 (3.7)	75.8	82.2	91.7	96.8	99.2	100	100

Table 4: Percentages of better or equal performance regarding the AIC-value of the two algorithms compared to Dißmann's algorithm for different dimensions d and sample sizes n based on 1000 data sets sampled from randomly generated R-vines (in brackets the percentages of equal performance are given).

time for dimensions larger than 20. Further, it seems that the sample size has no effect on which of the two algorithms performs better, since the percentages stay rather constant when the sample size changes from 400 to 3000.

Concerning Algorithm 2 the results are even more promising. In every scenario it is better or equal than Dißmann's algorithm in more than two thirds of the time. Especially in high dimensions ($d \geq 15$) it outperforms in more than 90% of simulations. Moreover, we note that a larger sample size helps to increase the advantage of Algorithm 2 even further with increasing percentages for n going from 400 to 3000, regardless of the considered dimension.

Regarding the question, how much better Algorithm 2 performs than Dißmann's algorithm we present in Figure 3 boxplots of the difference between the AIC-values per observation of the models found by Dißmann's and our second algorithm for $n = 1000$. Positive values imply a worse performance of the Dißmann algorithm.


Figure 3: Boxplots of the difference between the AIC-values per observation of the vines chosen by Dißmann's algorithm and Algorithm 2.

The boxplots affirm the results from Table 4. The differences between the AIC of Dißmann's algorithm and Algorithm 2 are mostly positive and become larger as the dimension increases. We further see that the magnitude of AIC improvements is much larger than that of the AIC losses. Since the values are shown per observations, we observe AIC improvements of up to 60,000 in 50 dimensions. Due to the scaling the values in the lower dimensions appear to be close to zero. However, we performed t-tests with the null hypothesis that the mean of the AIC differences is zero and could reject it in favor of Algorithm 2 with p-values smaller than 10^{-30} for all dimensions. The boxplots in the scenarios $n = 400$

and $n = 3000$ are very similar and can be seen in Appendix D.

In our simulation study the pair copula's Kendall's τ values were sampled from $[-1, 1]$. Since many real life data sets (especially in finance) consist of variables that are pairwise positively dependent we repeated our simulation study sampling only positive Kendall's τ values. An excerpt of the results is given in the first two rows of Table 8 in Appendix E. The results are quite similar with better performances in low and worse performance in high dimensions.

In the construction of the algorithms we noted that the pair-copula selection can be done with or without a prior independence test deciding whether the independence copula can be used. The results above were obtained without the use of independence test. In the last two lines of Table 8 in Appendix E we present the results in the positive dependence case when our as well as Dißmann's algorithms use the independence test at level $\beta = 0.05$. All the results are very similar to those without the independence test implying that the relative performances of the algorithms do not strongly depend on the use of the independence test. Finally, for both algorithms we chose the rank functions for scoring the Kendall's τ and p-values of the test for constant conditional correlations. This is justified in Appendix F.

4.4 Computational times and a faster version of Algorithm 2

We consider whether the better performance of our algorithm comes at a higher computational cost. The average computational times (based on 100 repetitions) of fitting a vine copula with the three competing algorithms depending on the dimension d are displayed for $n = 1000$ in Table 5.

					d				
Algorithm	3	4	5	7	10	15	20	30	50
Dißmann	0.42	0.92	1.47	3.14	6.63	15.6	24.6	67.8	155.3
Algorithm 1	0.41	0.83	1.42	2.94	6.42	15.2	23.5	65.1	157.7
Algorithm 2	0.49	1.20	2.37	6.65	20.08	72.9	154.6	731.2	3409.4

Table 5: Average computational times of the three algorithms for $n = 1000$ in different dimensions.

We see that the computational times of Dißmann's algorithm and Algorithm 1 are always very close to each other with Algorithm 1 being even faster most of the times. Algorithm 2 clearly is much slower than the other algorithms since for the determination of the optimal root node each possible pair-copula of the current tree level has to be fitted for the calculation of all the required p-values. This effect increases with the dimension making Algorithm 2 comparably fast in low dimensions, only 3 times slower in 10 dimensions, but roughly 20 times slower for $d = 50$.

A possible way to decrease the computational time of Algorithm 2 is given by the following adjustment: in a first step we restrict the allowed pair-copula families exclusively to the Gaussian copula and proceed as before to determine the estimated tree structure. In a second step all pair-copulas appearing in this tree structure are fitted again, now allowing for all pair-copula families. The idea is that the resulting vine copula models hopefully do not differ too much from the unadjusted fit, since the p-values of the CCC tests should not be very sensitive to the chosen families used to determine the pseudo observations and thus yield similar choices for the optimal root nodes. At the same time the computational effort is greatly reduced, since the most time consuming step of Algorithm 2 is the fitting of the pair-copulas to determine the p-values of the CCC tests. Table 6 displays the percentages

of computational times and AIC-values of the adjusted version of Algorithm 2 relative to the unadjusted version for dimensions $d = 5, 10, 30$ and sample sizes $n = 400, 1000, 3000$.

d	5			10			30		
n	400	1000	3000	400	1000	3000	400	1000	3000
% time	52.0	52.1	51.9	31.2	29.6	28.2	15.8	14.0	12.9
% AIC	98.5	97.7	97.1	98.4	98.7	97.4	99.0	98.2	98.2

Table 6: Percentages of computational times and AIC-values of the adjusted version of Algorithm 2 relative to the unadjusted version for dimensions $d = 5, 10, 30$ and sample sizes $n = 400, 1000, 3000$.

We see that the adjustment greatly reduces the computational time making the algorithm almost as fast as Dißmann's algorithm. Further, the loss in performance is negligibly small with relative performances ranging between 97 and 99 percent. Thus, we found a way to significantly improve the computational times of Algorithm 2 without losing too much performance.

5 Real data examples

In order to assess the performance of our proposed algorithms when confronted with real data, we revisit several data sets that have been used in the recent vine copula literature. For the description of each of the data sets we refer the reader to the respective references. The uranium data set, of which we already examined a three-dimensional subset in Section 3.1 was introduced in Cook and Johnson (1986). The bike data sets (Schallhorn et al., 2017) are daily and hourly records of bike rentals in Washington, D.C., together with local climate data (temperature, perceived temperature, humidity and wind speed). The hourly bike data set additionally contains the variable hour. The MAGIC (Major Atmospheric Gamma-ray Imaging Cherenkov) data has been analyzed in the context of classification by Nagler and Czado (2016). We only show the results for the subset with the classification hadron. Further, the Norwegian data is a financial data set of Norwegian and international market variables used in the context of truncated regular vines by Brechmann et al. (2012) and Killiches et al. (2016). Two higher dimensional financial data sets are given by the CDS data (Brechmann et al., 2013; Kraus and Czado, 2017) and the EuroStoxx 50 data (Brechmann and Czado, 2013; Killiches et al., 2016). Finally, in order to include another non-financial data set, we also consider a subset of the Concrete Compressive Strength Data Set (Yeh, 1998), containing all of its continuous variables, namely *Cement*, *Coarse aggregate*, *Fine aggregate* and *Concrete compressive strength*.

The dimensions of the data sets range between 3 and 52 with sample sizes between 655 and 17379. All of the data was transformed to the copula scale using adequate preprocessing methods where necessary (e.g. GARCH models for time series data as in Liu and Luger, 2009) and empirical probability integral transforms. Table 7 shows the AIC-values of the vine copulas fitted by the three competing algorithms. We further provide in brackets for each vine copula fitted by the three algorithms the number of pair-copulas for which the test for constant conditional correlations is rejected at the 5% level. So, e.g. in the case of the 3-dimensional uranium data set the CCC test is rejected with a p-value of 0.01 for the conditional copula modeled by Dißmann's algorithm (yielding a "1" in brackets) while Algorithm 2 chooses a structure with no rejection (implying the "0").

Similar to its three-dimensional subset the complete seven-dimensional uranium data set is best fitted by Algorithm 2. Out of the 15 conditional pair-copulas of the vine fitted by

Data	uranium	concrete	bike daily	bike hourly	uranium
d	3	4	5	6	7
n	655	1030	731	17379	655
Dißmann	-845.6 (1)	-522.9 (2)	-4423.5 (3)	-89859.6 (9)	-1759.9 (4)
Algorithm 1	-845.6 (1)	-522.9 (2)	-4423.5 (3)	-89842.4 (9)	-1756.4 (4)
Algorithm 2	$\mathbf{-857.8}$ (0)	$\mathbf{-551.6}$ (3)	$\mathbf{-4454.1}$ (3)	$\mathbf{-92040.5}$ (10)	$\mathbf{-1817.2}$(1)

Data	MAGIC	Norwegian	CDS	EuroStoxx
d	10	19	38	52
n	6688	1187	1371	985
Dißmann	-61359.4 (23)	-12906.4 (18)	$\mathbf{-40985.7}$ (31)	$\mathbf{-62377.8}$ (87)
Algorithm 1	$\mathbf{-61649.7}$(24)	$\mathbf{-12946.6}$ (10)	-40922.4 (25)	-62280.1 (65)
Algorithm 2	-58197.4 (17)	-12756.5 (12)	-40514.5 (32)	-62074.8 (68)

Table 7: AIC-values of the three structure selection methods applied to real data sets. In brackets the number of pair-copulas is given for which the CCC test is rejected at 5% level.

Dißmann's algorithm four are deemed non-simplified at 5% level. Algorithm 2 is able to reduce this number to one, thus yielding a better model fit. Regarding the four-dimensional concrete data set we see that Algorithm 2 is also able to improve the model fit. It may seem surprising, that compared to the other methods it increases the number of pair-copulas that violate the simplifying assumption. However, if we have a closer look at the p-values of the fitted pair-copulas, we see that the Dißmann fitted vine contains one pair-copula with p-value equal to zero which apparently severely impairs the model's likelihood. Similar observations can be made for the two bike data sets.
Considering the higher dimensional data sets we notice that Algorithm 2 looses its competitive advantage. The ten-dimensional MAGIC data as well as the 19-dimensional Norwegian data achieve their best fit using Algorithm 1. For the MAGIC data set it increases the overall non-simplified pair-copulas while reducing the number of strong violations (with p-value equal to zero) from 12 to 8. The good performance of Dißmann's algorithm for the high-dimensional finance data sets is not surprising. It is a stylized fact that these are modeled fairly well by t copulas (see for example Demarta and McNeil, 2005). Stöber et al. (2013) have shown that vine decompositions of t copulas are of the simplified form, independently of the vine's tree structure. Thus, the CCC test is rejected less often for vines modeling financial data. So for the 38-dimensional CDS data set, in the Dißmann vine only 31 of the 666 fitted pair-copulas violate the simplifying assumption, which is slightly less than the 5% that we would expect if the null hypothesis would be true for all pair-copulas. Thus, for structure selection purposes the p-value score is rather unimportant in this example. Consequently, Dißmann's heuristic of modeling the strongest dependencies in the lower trees yields the best results and the restriction to C-vines imposed by Algorithm 2 is apparently too severe. Similar arguments hold for the 52-dimensional EuroStoxx data set, where only 6.8% of the pair-copulas fitted by Dißmann's algorithm violate the simplifying assumption. Regarding the pair-copula families selected by the three algorithms, we observe that they do not differ considerably between the algorithms. Not surprisingly, most of the pair-copula families of the financial data sets are t-copulas, with some Tawn, Frank and BB1 copulas in the mix. The concrete data was modeled mainly by Tawn, Frank and Gauss copulas. The same copula families were chosen for the two bike data sets, with some additional BB8 copulas and rotations thereof. For the uranium data set, Frank, Joe, t and Tawn copulas were favored, while the most commonly selected families of the MAGIC data were Tawn,

BB8 and t copulas. So we see that in these examples many non-Gaussian copulas are selected such that Gaussian copulas would not provide satisfying model fits.

All in all, we have seen that the algorithms proposed in this paper are able to improve the model fit compared to Dißmann's algorithm, especially when the Dißmann vine exhibits significant non-simplifiedness as assessed by the CCC test.

6 Conclusion

We propose two new algorithms for the sequential selection of the tree structure of a vine copula model. We extend currently existing methods by incorporating tests which gage the validity of the simplifying assumption for every pair-copula. This results in vine copula models which frequently have a better model fit than the benchmark given by Dißmann's algorithm.

We have revisited the three-dimensional uranium data set which is famous for its non-simplified nature when considering its vine decomposition using Dißmann's algorithm. However, we found out that using the vine decomposition where the conditioning in the second tree is done with respect to cobalt, the resulting vine is of the simplified form and provides a better model fit.

In the real data application we have seen that our proposed algorithms work especially well, when the vine fitted by Dißmann's algorithm contains many pair-copulas violating the simplifying assumption. Thus, from a practitioner's point of view, after fitting a d-dimensional vine copula using Dißmann's algorithm one should always count the number of pair-copulas for which the test of constant conditional correlations is rejected at some level β. If this number is considerably larger than $\beta(d-1)(d-2)/2$ (the expected number of rejections), one should refit the model using our proposed algorithms, accepting larger computational times as a trade-off for a likely better model fit.

Appendix

A Pair-copula families implemented in `VineCopula`

The pair-copula families currently implemented in the `VineCopula` package are Gaussian, t*, Clayton, Gumbel, Frank, Joe, BB1*, BB6*, BB7*, BB8* and Tawn* with respective rotations. The stars indicate two-parametric families.

B Simulation of the second parameter for two-parametric copula families

t copula   The degrees of freedom of the t copula are sampled from $3 + G$, where $G \sim Gamma(3, 3)$. This implies an expected value of 4 and the 95% confidence interval $[3.2, 5.4]$.

BB1, BB6, BB7   The second parameters of the BB1, BB6 and BB7 copulas are sampled from $1 + 3B$, where $B \sim Beta(4, 2)$. This implies an expected value of 3 and the 95% confidence interval $[1.84, 3.85]$. When the dependence is negative, the second parameter is multiplied with -1.

BB8, Tawn   The second parameters of the BB8 and Tawn copulas are sampled from a $Beta(4, 2)$ distribution. This implies an expected value of $2/3$ and the 95% confidence interval $[0.28, 0.95]$. When the dependence is negative, the second parameter of the BB8 copula is multiplied with -1.
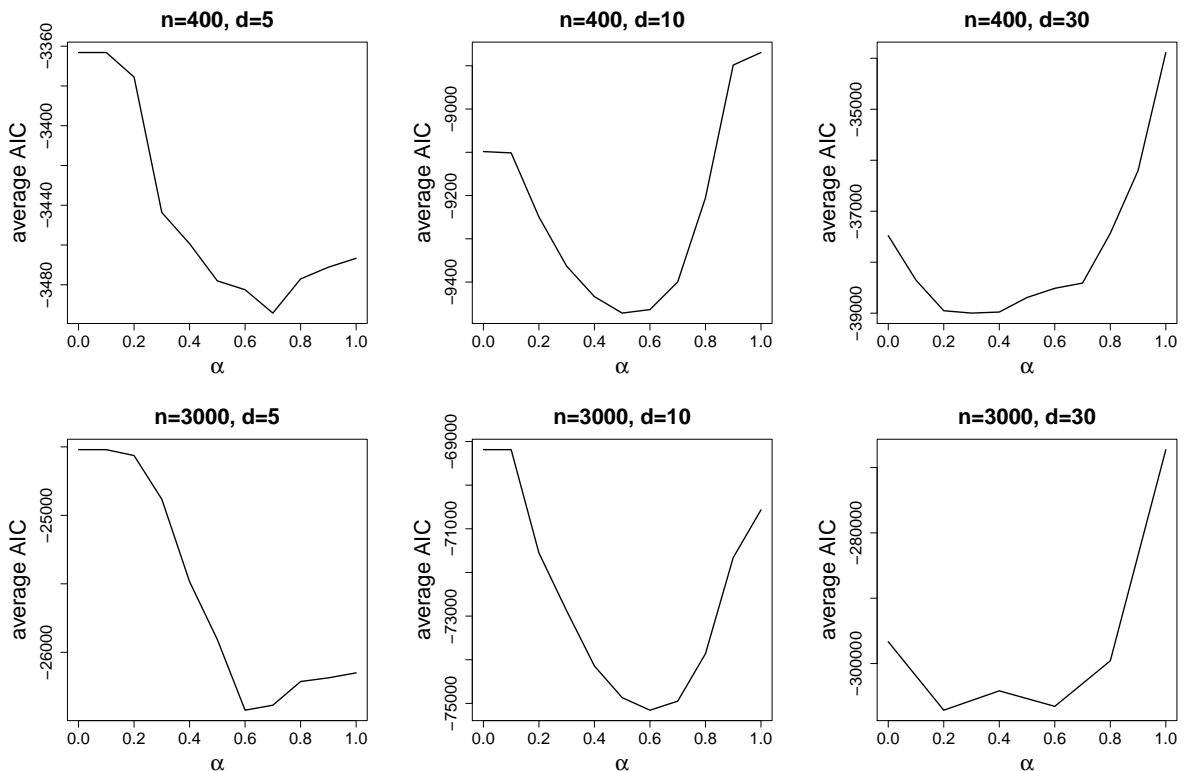
C Choice of weighting factor α

Figure 4: AIC-values of vine copulas fitted by Algorithm 2 depending on the weighting factor α in the settings $n = 400$ (top row) and $n = 3000$ (bottom row) for $d = 5, 10, 30$, averaged over 1000 repetitions.

D   Differences of AIC-values

Figure 5: Boxplots of the difference between the AIC-values per observation of the vines chosen by Dißmann's algorithm and Algorithm 2 for $n = 400$ (left panel) and $n = 3000$ (right panel).

E   Simulation study results for positive dependence and independence tests

Ind. test	Algorithm	d			
		3	5	10	30
Without	Algorithm 1	100 (100)	73.1 (53.4)	37.0 (3.2)	18.8
	Algorithm 2	94.1 (22.9)	83.1	86.1	91.9
With	Algorithm 1	100 (100)	73.6 (54.0)	33.0 (3.2)	16.8
	Algorithm 2	94.4 (22.9)	82.9	85.8	92.9

Table 8: Percentages of better or equal performance regarding the AIC-value of the two algorithms compared to Dißmann's algorithm for $n = 1000$ and random vine copulas with only positive dependence (in brackets the percentages of equal performance are given). In the first two rows pair-copula no independence tests were performed and in the last two rows they were performed with level $\beta = 0.05$.

F   Choice of p-value transformation function r

r	d			
	3	5	10	30
rank	91.1 (51.5)	73.6	88.7	99.6
identity	89.9 (48.5)	71.1	87.5	99.5
logarithm	89.9 (48.5)	72.4	77.0	94.8

Table 9: Percentages of better or equal performance regarding the AIC-value of Algorithm 2 compared to Dißmann's algorithm depending on the transformation function r for $n = 1000$ and $d = 3, 5, 10, 30$ (in brackets the percentages of equal performance are given).
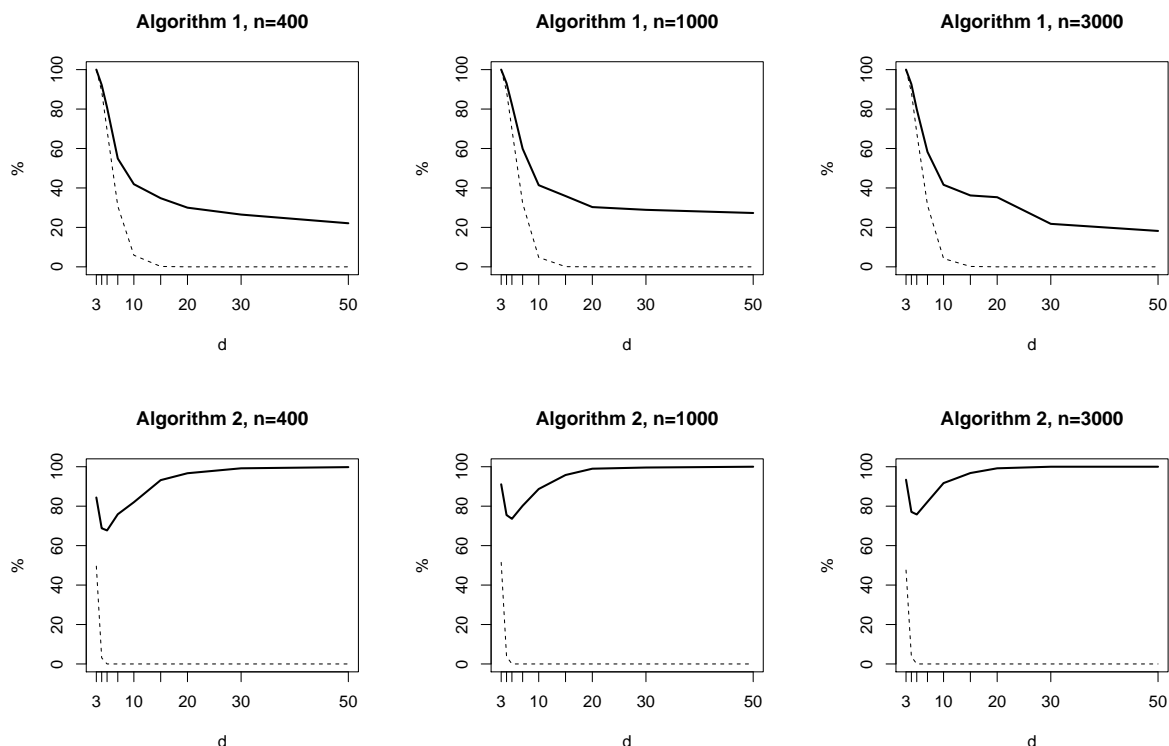
G Plots of simulation study results

Figure 6: Percentages of better or equal performance regarding the AIC-value of the two algorithms compared to Dißmann's algorithm for different dimensions d and sample sizes n based on 1000 data sets sampled from randomly generated R-vines (the dashed lines represent the percentages of equal performance).

References

Aas, K. (2016), "Pair-Copula Constructions for Financial Applications: A Review," *Econometrics*, 4.

Aas, K., Czado, C., Frigessi, A., and Bakken, H. (2009), "Pair-copula constructions of multiple dependence," *Insurance: Mathematics and economics*, 44, 182–198.

Acar, E. F., Genest, C., and Nešlehová, J. (2012), "Beyond simplified pair-copula constructions," *Journal of Multivariate Analysis*, 110, 74–90.

Akaike, H. (1998), "Information theory and an extension of the maximum likelihood principle," in *Selected Papers of Hirotugu Akaike*, Springer, pp. 199–213.

Almeida, C., Czado, C., and Manner, H. (2016), "Modeling high-dimensional time-varying dependence using dynamic D-vine models," *Applied Stochastic Models in Business and Industry*, 32, 621–638.

Barthel, N., Geerdens, C., Killiches, M., Janssen, P., and Czado, C. (2016), "Vine copula based inference of multivariate event time data," *arXiv preprint arXiv:1603.01476*.

Bedford, T. and Cooke, R. M. (2002), "Vines: A new graphical model for dependent random variables," *Annals of Statistics*, 30, 1031–1068.

Brechmann, E. C. and Czado, C. (2013), "Risk management with high-dimensional vine copulas: An analysis of the Euro Stoxx 50," *Statistics & Risk Modeling*, 30, 307–342.

Brechmann, E. C., Czado, C., and Aas, K. (2012), "Truncated regular vines in high dimensions with application to financial data," *Canadian Journal of Statistics*, 40, 68–85.

Brechmann, E. C., Hendrich, K., and Czado, C. (2013), "Conditional copula simulation for systemic risk stress testing," *Insurance: Mathematics and Economics*, 53, 722–732.

Cook, R. and Johnson, M. (1986), "Generalized burr-pareto-logistic distributions with applications to a uranium exploration data set." *Technometrics*, 28, 123–131.

Cooke, R. M., Joe, H., and Chang, B. (2015), "Vine Regression," *Resources for the Future Discussion Paper*, 15–52.

Cormen, T. H., Leiserson, C. E., Rivest, R. L., and Stein, C. (2001), *Introduction to algorithms*, vol. 6, MIT press Cambridge.

Czado, C. (2010), "Pair-copula constructions of multivariate copulas," in *Copula theory and its applications*, Springer, pp. 93–109.

Czado, C., Jeske, S., and Hofmann, M. (2013), "Selection strategies for regular vine copulae," *Journal de la Société Française de Statistique*, 154, 174–191.

Demarta, S. and McNeil, A. J. (2005), "The t copula and related copulas," *International Statistical Review/Revue Internationale de Statistique*, 111–129.

Diestel, R. (2005), *Graph Theory (Graduate Texts in Mathematics)*, Springer.

Dißmann, J., Brechmann, E. C., Czado, C., and Kurowicka, D. (2013), "Selecting and estimating regular vine copulae and application to financial returns," *Computational Statistics & Data Analysis*, 59, 52–69.

Erhardt, T. M. and Czado, C. (2015), "Standardized drought indices: A novel uni-and multivariate approach," *arXiv preprint arXiv:1508.06476*.

Genest, C. and Favre, A.-C. (2007), "Everything you always wanted to know about copula modeling but were afraid to ask," *Journal of hydrologic engineering*, 12, 347–368.

Haff, I. H., Aas, K., and Frigessi, A. (2010), "On the simplified pair-copula construction – Simply useful or too simplistic?" *Journal of Multivariate Analysis*, 101, 1296–1310.

Hobæk Haff, I., Frigessi, A., and Maraun, D. (2015), "How well do regional climate models simulate the spatial dependence of precipitation? An application of pair-copula constructions," *Journal of Geophysical Research: Atmospheres*, 120, 2624–2646.

Joe, H. (1997), *Multivariate models and multivariate dependence concepts*, CRC Press.

Joe, H., Cooke, R. M., and Kurowicka, D. (2011), "Regular vines: generation algorithm and number of equivalence classes," *Dependence Modeling: Vine Copula Handbook*, 219–231.

Killiches, M. and Czado, C. (2015), "Block-Maxima of Vines," in *Extreme Value Modelling and Risk Analysis: Methods and Applications*, eds. Dey, D. and Yan, J., Boca Raton, FL: Chapman & Hall/CRC Press, pp. 109–130.

Killiches, M., Kraus, D., and Czado, C. (2016), "Using model distances to investigate the simplifying assumption, goodness-of-fit and truncation levels for vine copulas," *arXiv preprint arXiv:1610.08795*.

— (2017), "Examination and visualisation of the simplifying assumption for vine copulas in three dimensions," *Australian & New Zealand Journal of Statistics*, 59, 95–117.

Kraus, D. and Czado, C. (2017), "D-vine copula based quantile regression," *Computational Statistics and Data Analysis*, 110C, 1–18.

Kurowicka, D. (2011), "Optimal truncation of vines," *Dependence Modeling: Vine Copula Handbook*, 233.

Kurz, M. S. (2017), *pacotest: Testing for Partial Copulas and the Simplifying Assumption in Vine Copulas*, version 0.2.

Kurz, M. S. and Spanhel, F. (2017), "Testing the simplifying assumption in high-dimensional vine copulas," *Unpublished working paper*.

Liu, Y. and Luger, R. (2009), "Efficient estimation of copula-GARCH models," *Computational Statistics & Data Analysis*, 53, 2284–2297.

Maya, L., Albeiro, R., Gomez-Gonzalez, J. E., and Melo Velandia, L. F. (2015), "Latin American exchange rate dependencies: A regular vine copula approach," *Contemporary Economic Policy*, 33, 535–549.

Morales-Nápoles, O. (2011), "Counting vines," in *Dependence Modeling: Vine Copula Handbook*, eds. Kurowicka, D. and Joe, H., Singapore, SG: World Scientific, chap. 9, pp. 189–218.

Nagler, T. and Czado, C. (2016), "Evading the curse of dimensionality in nonparametric density estimation with simplified vine copulas," *Journal of Multivariate Analysis*, 151, 69 – 89.

Nelsen, R. B. (2007), *An introduction to copulas*, Springer Science & Business Media.

Pereira, G., Veiga, Á., Erhardt, T., and Czado, C. (2016), "Spatial R-vine copula for streamflow scenario simulation," in *Power Systems Computation Conference (PSCC), 2016*, IEEE, pp. 1–7.

Schallhorn, N., Kraus, D., Nagler, T., and Czado, C. (2017), "Two approaches for discrete D-vine quantile regression," *Unpublished working paper*.

Schellhase, C. and Spanhel, F. (2017), "Estimating Non-Simplified Vine Copulas Using Penalized Splines," *Statistics and Computing*.

Schepsmeier, U., Stöber, J., Brechmann, E. C., Graeler, B., Nagler, T., and Erhardt, T. (2016), *VineCopula: Statistical Inference of Vine Copulas*, version 2.0.5.

Spanhel, F. and Kurz, M. S. (2015), "Simplified vine copula models: Approximations based on the simplifying assumption." *arXiv preprint arXiv:1510.06971*.

Stöber, J., Joe, H., and Czado, C. (2013), "Simplified pair copula constructions—limitations and extensions," *Journal of Multivariate Analysis*, 119, 101–118.

Vatter, T. and Nagler, T. (2016), "Generalized additive models for pair-copula constructions," *arXiv preprint arXiv:1608.01593*.

Yeh, I.-C. (1998), "Modeling of strength of high-performance concrete using artificial neural networks," *Cement and Concrete research*, 28, 1797–1808.